

CprE 419 Lab 3: Graph Processing using MapReduce

Department of Electrical and Computer Engineering
Iowa State University
Spring 2013

Purpose

In this lab, you will use Hadoop MapReduce for analyzing large graphs. A graph is a fundamental structure used for modeling relationships between entities, for example, hyperlinks between webpages, or friendship between people in a social network. At the end of this lab, you will know how to:

- Process a large graph that is presented as a set of edges
- Compute local properties of graphs, such as edge neighborhoods and triangles

Submission

Create a zip (or tar) archive with the following and hand it in through blackboard.

- A write-up answering questions for each experiment in the lab. For output, cut and paste results from your terminal and summarize when necessary.
- Commented Code for your program. Include all source files needed for compilation.

Experiments

Our dataset is the U.S. patent citation data, which is maintained by the National Bureau of Economic Research. In the graph that is considered, the vertex set is the set of all patents issued between 1975 and 1999, for a total of nearly 4 million patents. For each citation, say, from patent A to patent B, there is an edge from vertex representing A to the vertex representing B, in the citation graph. Hence this citation graph is a directed graph.

The above graph has been uploaded to the server at the location: `"/datasets/Lab3/Graph"`. This graph is a subset of the original, constructed by the Stanford Network Analysis Project. More information about the data can be obtained from its source:

<http://snap.stanford.edu/data/cit-Patents.html>

The graph is in the form of an edge list. Every line of the file has information about a single edge. A line contains information in the format `<from vertex> <to vertex>`, which means that patent `<from vertex>` has a citation to patent `<to vertex>`.

Experiment 1 (40 pts)

The first task is to find out significant patents, defined as follows. We say that there is a two-hop citation from patent X to patent Y if there is a patent Z such that X cites a patent Z and Z cites Y. We say that there is a one-hop citation from X to Y if X cites Y directly. For the purposes of this experiment, we define the significance of a patent X to equal the number of distinct patents Y such that there is either a one-hop citation or a two-hop citation from Y to X. Your task is the following: *Write a MapReduce program to extract the ten patents with the largest significance.* If there is a tie in choosing the winners, then they can be broken arbitrarily.

You should output the top ten patents and their significance in the directory
"/users/<User ID>/Lab3/exp1/"

Experiment 2 (40 pts)

For the next experiment consider the same patent graph as the input, but convert it into an undirected graph by deleting the direction on an edge. Thus each edge in the input file represents an undirected edge between the two vertices.

A triangle is a set of three vertices such that all three pairs of vertices are connected to each other. For example vertices {4, 7, 9} form a triangle in a graph if and only if the graph has the following edges: {4, 7}, {4, 9}, {7, 9}. The number of triangles in a graph is an important metric of a graph that has applications in several domains including social network analysis. Note that a single vertex can participate in multiple triangles in the graph.

Write a MapReduce program to compute the total number of triangles in the input undirected graph. You should output all your data under: "/users/<User ID>/Lab3/exp2/"